

# Pthread: Mutex

- ❏ Provides synchronization to a shared resource
- ❏ A typical sequence in the use of a mutex is as follows:
  - ❏ Create and initialize a mutex variable
  - ❏ Several threads attempt to lock the mutex
  - ❏ Only one succeeds and that thread owns the mutex
  - ❏ The owner thread performs some set of actions
  - ❏ The owner unlocks the mutex
  - ❏ Another thread acquires the mutex
  - ❏ ... ..
  - ❏ Finally the mutex is destroyed

# POSIX Thread Library

## # include <pthread.h>

✧ int **pthread\_mutex\_init** (pthread\_mutex\_t \*mutex,  
const pthread\_mutexattr\_t \*attr)

✧ initializes a *mutex* with the specified attributes

✧ 1<sup>st</sup> argument: the address of the newly created *mutex*

✧ 2<sup>nd</sup> argument: the address to the variable containing the mutex attributes object

✧ To use the default attributes, use NULL as the second argument

# POSIX Thread Library

```
# include <pthread.h>
```

✧ int **pthread\_mutex\_destroy** (pthread\_mutex\_t \*mutex)

✧ Free a mutex object that is no longer needed.

# POSIX Thread Library

## # include <pthread.h>

✱ int **pthread\_mutex\_lock** (pthread\_mutex\_t \*mutex)

- ✱ acquires ownership of the *mutex* specified
- ✱ If the specified *mutex* is currently locked, the calling thread is blocked until the *mutex* is available
- ✱ 1<sup>st</sup> argument: the address of the *mutex* to be locked

📖 int **pthread\_mutex\_unlock** (pthread\_mutex\_t \*mutex)

- 📖 unlocks the *mutex* specified
- 📖 1<sup>st</sup> argument: the address of the *mutex* to be unlocked

📖 Important: make sure all threads, that need access to the shared data, use mutex. Otherwise, the shared data could still be corrupted.

# Pthread: Condition variable

- ❏ A mutex makes other threads to wait while the thread holding the mutex executes code in a critical section.
- ❏ A condition variable is typically used by a thread (**that has already acquired a mutex**) to make itself wait (**and temporarily release the mutex**) until signaled by another thread.



# POSIX Thread Library

## # include <pthread.h>

✧ int **pthread\_cond\_wait** (pthread\_cond\_t \*cond,  
pthread\_mutex\_t \*mutex)

✧ blocks the calling thread on the condition variable *cond*

✧ 1<sup>st</sup> argument: the address of the condition variable to wait on

✧ 2<sup>nd</sup> argument: the address of the mutex associated with the condition variable

✧ when **pthread\_cond\_wait()** is called, the calling thread must have the associated *mutex* locked

✧ the **pthread\_cond\_wait()** function unlocks the associated *mutex* and blocks on the condition variable (waiting for another thread to signal the condition)

# POSIX Thread Library

## # include <pthread.h>

- ✧ int **pthread\_cond\_signal** (pthread\_cond\_t \*cond)
  - ✧ wakes up one thread that is waiting on the condition variable
  - ✧ 1<sup>st</sup> argument: the address of the condition variable
  - ✧ the thread that is just waked up automatically requests for the mutex that it unlocked when calling *pthread\_cond\_wait()*; it can resume its execution only after it has re-acquired the *mutex*



# Example: Implementing RW-Monitor

```
/*Shared variables*/
```

```
int readercount;
```

```
int busy;
```

```
pthread_mutex_t mutex;
```

```
pthread_cond_t OKtoread, OKtowrite;
```

```
int QL_OKtoread;
```

**Monitor** Reader-Writer

begin

    readercount: integer

    busy: boolean

    OKtoread, OKtowrite: condition

# Example: Implementing RW-Monitor

```
/*Initialization*/
```

```
void init()
```

```
{
```

```
pthread_mutex_init(&mutex, NULL);
```

```
pthread_cond_init(&OKtoread, NULL);
```

```
pthread_cond_init(&OKtowrite, NULL);
```

```
readercount=0;
```

```
busy=0;
```

```
QL_OKtoread=0;
```

```
}
```

```
begin /* initialization of local  
data */
```

```
    readercount := 0;
```

```
    busy := false;
```

```
end
```

# Example: Implementing RW-Monitor

```
/*Procedure startread*/  
void startread()  
{  
    pthread_mutex_lock(&mutex);  
    QL_OKtoread++;  
    while(busy)  
        pthread_cond_wait(  
            &OKtoread, &mutex);  
    QL_OKtoread--;  
    readercount++;  
    pthread_cond_signal(&OKtoread);  
    pthread_mutex_unlock(&mutex);  
}
```

```
procedure startread  
begin  
    If busy then OKtoread.wait;  
    readercount := readercount + 1;  
    OKtoread.signal  
end startread
```

# Example: Implementing RW-Monitor

```
/*Procedure endread*/  
void endread()  
{  
    pthread_mutex_lock(&mutex);  
    readercount--;  
    if(readercount==0)  
        pthread_cond_signal(&OKtowrite);  
    pthread_mutex_unlock(&mutex);  
}
```

```
procedure endread;  
begin  
    readercount := readercount-1;  
    If readercount = 0 then  
        OKtowrite.signal;  
end endread
```

# Example: Implementing RW-Monitor

```
/*Procedure startwrite*/  
void startwrite()  
{  
    pthread_mutex_lock(&mutex);  
    while(busy || readercount!=0)  
        pthread_cond_wait(&OKtowrite,  
            &mutex);  
    busy=1;  
    pthread_mutex_unlock(&mutex);  
}
```

```
procedure startwrite;  
begin  
    If busy or readercount <>0 then  
        OKtowrite.wait;  
    busy := true;  
end startwrite
```

# Example: Implementing RW-Monitor

```
/*Procedure endwrite*/  
void endwrite()  
{  
    pthread_mutex_lock(&mutex);  
    busy=0;  
    if(QL_OKtoread>0)  
        pthread_cond_signal(&OKtoread);  
    else  
        pthread_cond_signal(&OKtowrite);  
    pthread_mutex_unlock(&mutex);  
}
```

```
procedure endwrite;  
begin  
    busy := false  
    if OKtoread.queue then  
        OKtoread.signal  
    else OKtowrite.signal  
end endwrite
```

# Example: Using RW-Monitor

```
/*reader*/  
void *reader(void *x)  
{  
    startread();  
    //code to read  
    endread();  
}
```

```
/*writer*/  
void *writer(void *x)  
{  
    startwrite();  
    //code to write  
    endwrite();  
}
```

# Deadlock (I)

September 25, 2017



# The Deadlock Problem

❏ A set of blocked processes each holds an instance of resource and waits to acquire an instance of resource held by another process

❏ Example

❏ System has 2 disk drives

❏  $P_1$  and  $P_2$  each hold one disk drive and each needs another one

❏ Example

❏ semaphores  $A$  and  $B$ , initialized to 1

$P_0$	$P_1$
wait (A);	wait(B)
wait (B);	wait(A)

# Deadlock Characterization

Deadlock arises → 4 conditions hold simultaneously  
(necessary conditions of deadlock)

- ❏ **Mutual exclusion:** an instance of resource can only be used by a process at a time
- ❏ **Hold and wait:** a process holding at least one instance of resource is waiting to acquire additional resources held by other processes
- ❏ **No preemption:** resource can be released only voluntarily by the process holding it, after that process has completed its task
- ❏ **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .

# System Model

- Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
- Each process utilizes a resource in the following order:  
**request  $\rightarrow$  use  $\rightarrow$  release**

# Resource-Allocation Graph

A set of vertices  $V$  and a set of edges  $E$ .

☐  $V$  is partitioned into two types:

☐  $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system


☐  $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system

☐ **request edge** – directed edge  $P_i \rightarrow R_j$


☐ **assignment edge** – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph (Cont.)


- Process
- Resource Type with 4 instances




- $P_i$  requests an instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$



# Example of a Resource Allocation Graph




# Resource Allocation Graph With A Deadlock



There is one or more cycles in the graph

# Graph With A Cycle But No Deadlock





# Basic Facts

- ❏ If graph contains no cycles  $\Rightarrow$  no deadlock
- ❏ If graph contains a cycle  $\Rightarrow$ 
  - ❏ if only one instance per resource type, then deadlock
  - ❏ if several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- ❏ Ensure that the system will *never* enter a deadlock state
- ❏ Allow the system to enter a deadlock state and then recover
- ❏ Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX

# Deadlock Prevention

Restrain the ways that request can be made such that at least one of the deadlock necessary conditions fails

- ❏ **Mutual Exclusion** → **Make each instance of resource simultaneously sharing** – often impossible
- ❏ **Hold and Wait** → **Request all at once; or, Give up all before request** – must guarantee that whenever a process requests a resource, it does not hold any other resources (i.e., it should release all resources that it currently holds)
  - ❏ Low resource utilization; starvation possible

# Deadlock Prevention

- ❏ **No Preemption** → **Allow resource instance to be re-assignable**
  - ❏ Often impossible
- ❏ **Circular Wait** → **Prevent Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration

# Deadlock due to circular wait

P1 & P2 share semaphores A & B & C

P1:

wait(A)

...segment 1...

wait(B)

...segment 2...

wait(C)

...segment 3 ...

signal(C)

signal(B)

signal(A)

P2:

wait(B)

...segment 4...

wait(C)

...segment 5...

wait(A)

...segment 6 ...

signal(A)

signal(C)

signal(B)

# Prevent circular wait

P1 & P2 share semaphores A & B & C; impose order:  $A < B < C$

P1:

wait(A)

...segment 1...

wait(B)

...segment 2...

wait(C)

...segment 3 ...

signal(C)

signal(B)

signal(A)

P2:

wait(A)

wait(B)

...segment 4...

wait(C)

...segment 5...

...segment 6 ...

signal(A)

signal(C)

signal(B)