

Deadlock (III)

October 4, 2017

Banker's Algorithm

- ❏ Multiple instances
- ❏ Each process must claim maximum use in advance
- ❏ When a process requests a resource it may have to wait
- ❏ When a process gets all its resources it must return them in a finite amount of time

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- 🖥️ **Available:** Vector of length m . If $available[j] = k$, there are k instances of resource type R_j available
- 🖥️ **Max:** $n \times m$ matrix. If $Max[i,j] = k$, then process P_i may request at most k instances of resource type R_j
- 🖥️ **Allocation:** $n \times m$ matrix. If $Allocation[i,j] = k$ then P_i is currently allocated k instances of R_j
- 🖥️ **Need:** $n \times m$ matrix. If $Need[i,j] = k$, then P_i may need k more instances of R_j to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

Safety Algorithm

1. Let $Work$ and $Finish$ be vectors of length m and n , respectively.

Initialize:

$$Work = Available$$

$$Finish[i] = false, \text{ for } i = 0, 1, \dots, n-1$$

2. Find an i such that both:

(a) $Finish[i] = false$

(b) $Need_i \leq Work$

If no such i exists, go to step 4

3. $Work = Work + Allocation_i$

$$Finish[i] = true$$

go to step 2

4. If $Finish[i] == true$ for all i , then the system is in a safe state

Resource-Request Algorithm for Process P_i

$Request$ = request vector for process P_i . If $Request_i[j] = k$ then process P_i wants k instances of resource type R_j

1. If $Request_i \leq Need_i$ go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
2. If $Request_i \leq Available$, go to step 3. Otherwise P_i must wait, since resources are not available
3. Pretend to allocate requested resources to P_i by modifying the state as follows:

$$Available = Available - Request;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If *safe* \Rightarrow the resources are allocated to P_i
- If *unsafe* $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Example of Banker's Algorithm

5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances)

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

Example (Cont.)

■ The content of the matrix *Need* is defined to be $Max - Allocation$

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
P_0	7	4	3
P_1	1	2	2
P_2	6	0	0
P_3	0	1	1
P_4	4	3	1

■ The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria

Example: P_1 Request (1,0,2)

☐ Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow \text{true}$)




	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

☐ Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement

☐ Can request for (3,3,0) by P_4 be granted?

☐ Can request for (0,2,0) by P_0 be granted?


Deadlock Detection

-  Allow system to enter deadlock state
-  Detection algorithm
-  Recovery scheme


Single Instance of Each Resource Type

- ❏ Maintain *wait-for* graph
 - ❏ Nodes are processes
 - ❏ $P_i \rightarrow P_j$ if P_i is waiting for P_j
- ❏ Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock
- ❏ An algorithm to detect a cycle in a graph requires an order of n^2 operations, where n is the number of vertices in the graph

Resource-Allocation Graph and Wait-for Graph






Resource-Allocation Graph



Corresponding wait-for graph

Several Instances of a Resource Type

-  **Available:** A vector of length m indicates the number of available resources of each type.
-  **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process.
-  **Request:** An $n \times m$ matrix indicates the current request of each process. If $Request[i][j] = k$, then process P_i is requesting k more instances of resource type R_j .

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length m and n , respectively
Initialize:
 - (a) $Work = Available$
 - (b) For $i = 1, 2, \dots, n$, if $Allocation_i \neq 0$, then
 $Finish[i] = false$; otherwise, $Finish[i] = true$
2. Find an index i such that both:
 - (a) $Finish[i] == false$
 - (b) $Request_i \leq Work$

If no such i exists, go to step 4

Detection Algorithm (Cont.)

3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2

4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state. Moreover, if $Finish[i] == false$, then P_i is deadlocked

Algorithm requires an order of $O(m \times n^2)$ operations to detect whether the system is in deadlocked state

Example of Detection Algorithm

Five processes P_0 through P_4 ; three resource types
A (7 instances), B (2 instances), and C (6 instances)

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

Execution sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i . So, the system is not deadlocked.

Example of Detection Algorithm

Five processes P_0 through P_4 ; three resource types
A (7 instances), B (2 instances), and C (6 instances)

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	<i>A B C</i>	<i>A B C</i>	<i>A B C</i>
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 1	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

Deadlock exists, consisting of processes P_1 , P_2 , P_3 , and P_4

Detection-Algorithm Usage

- ❏ When, and how often, to invoke detection?
 - ❏ Whenever a resource request is made
 - ❏ Maybe too frequent
 - ❏ If the maximum resource demand of every process is known, this can prevent deadlock
 - ❏ Whenever a resource request cannot be satisfied
 - ❏ Can identify the process which “finally” causes deadlock
 - ❏ Every certain time interval
 - ❏ There may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock

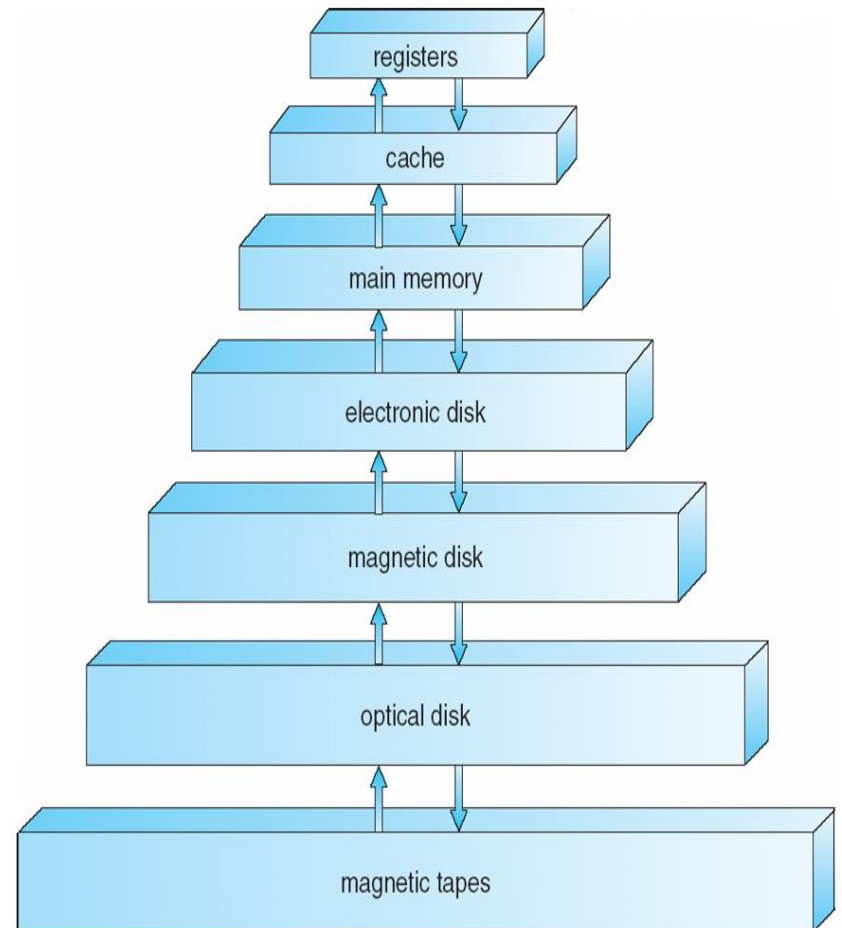
Recovery from Deadlock: Process Termination

- ❏ Abort all deadlocked processes
- ❏ Abort one process at a time until the deadlock cycle is eliminated
 - ❏ In which order should we choose to abort?
 - ❏ Priority of the process
 - ❏ How long process has computed, and how much longer to completion
 - ❏ Resources the process has used
 - ❏ Resources process needs to complete
 - ❏ Is process interactive or batch?

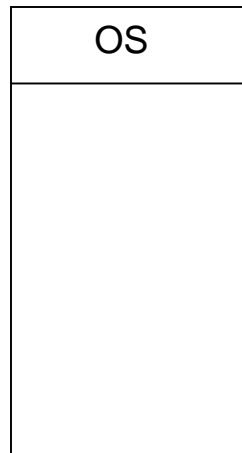
Storage Hierarchy

❑ Main memory, cache and registers are the only storages that CPU can access directly

❑ Program must be brought (from disk) into memory and placed within a process image for it to be run




- ❏ Main memory is usually divided into two partitions:
 - ❏ Resident operating system, usually held in low memory with interrupt vector
 - ❏ User processes then held in high memory



Logical Addresses

Physical Addresses

Used in code,
by CPU




Used by MMU
(memory
management
unit)

Main
Memory


Contiguous Allocation

- ❏ In user memory space, each process is stored in a contiguous region (block).
- ❏ Hole – block of un-occupied contiguous memory space
- ❏ At the beginning, there is a single hole in the memory: the whole space for user processes








Contiguous Allocation

- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Operating system maintains information about:
 - a) allocated partitions b) free partitions (holes)
- Holes of various sizes may be generated later on and are scattered throughout memory



Contiguous Allocation Policies

How to satisfy a request of size n from a list of free holes

-  **First-fit:** Allocate the *first* hole that is big enough
-  **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 -  Produces the smallest leftover hole
-  **Worst-fit:** Allocate the *largest* hole; must also search entire list
 -  Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization



Problem: Fragmentation

- ❏ **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- ❏ Reduce external fragmentation by **compaction**
 - ❏ Shuffle memory contents to place all free memory together in one large block
 - ❏ Compaction is possible *only* if relocation is dynamic, and is done at execution time

Paging – Memory management

Strategy adopted by modern OSes

Objective:

-  Logical address space of a process remains contiguous but the physical address space of it needs not be contiguous
-  Process is allocated physical memory whenever the latter is available

Paging: Key Ideas

- ❏ Divide physical memory (user memory part) into fixed-sized blocks called **frames** (size is power of 2, between 512 bytes and 8,192 bytes)
- ❏ Divide logical memory space of a process into blocks of same size called **pages**
- ❏ Pages are mapped to frames one-by-one; process-specific **page table** records the mapping and facilitates the logical to physical address translation
- ❏ OS keeps track of free frames and allocates frames to new/swap-in processes

Paging Model and Page Table

