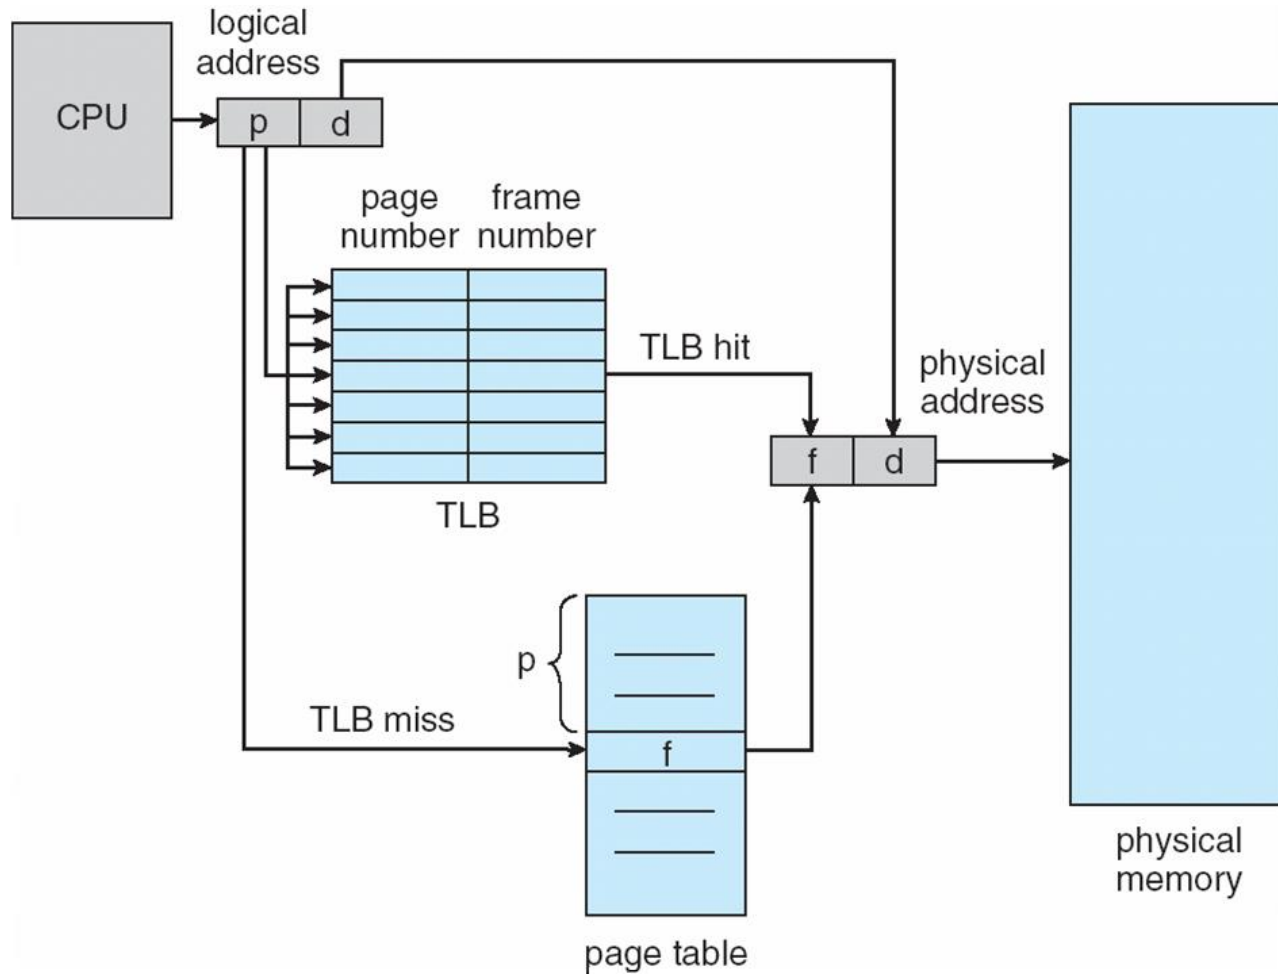


Main Memory III

October 6, 2017

Paging Hardware With TLB



TLB Hardware is shared by multiple processes

❏ Problem: when a process is swapped out and a new process is swapped in, the content of the TLB becomes outdated

❏ Solutions:

❏ flush the TLB when process switch; or

❏ store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process

Implementation of Page Table

- ❏ Page table is kept in main memory (kernel space)
- ❏ **Page-table base register (PTBR)** points to the page table
- ❏ **Page-table length register (PTLR)** indicates size of the page table
- ❏ Memory protection implemented by
 - ❏ PTLR specifies the length of page table (the number of pages for a process)
 - ❏ If the page number of an address \geq the value of PTLR, the logical address is invalid

Shared Pages

Shared code

-  One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

Private code and data

-  Each process keeps a separate copy of the code and data

Structure of the Page Table

❏ Number of logical pages could be very large →

❏ Hierarchical Paging

❏ Inverted Page Tables

❏ Hashed Page Tables

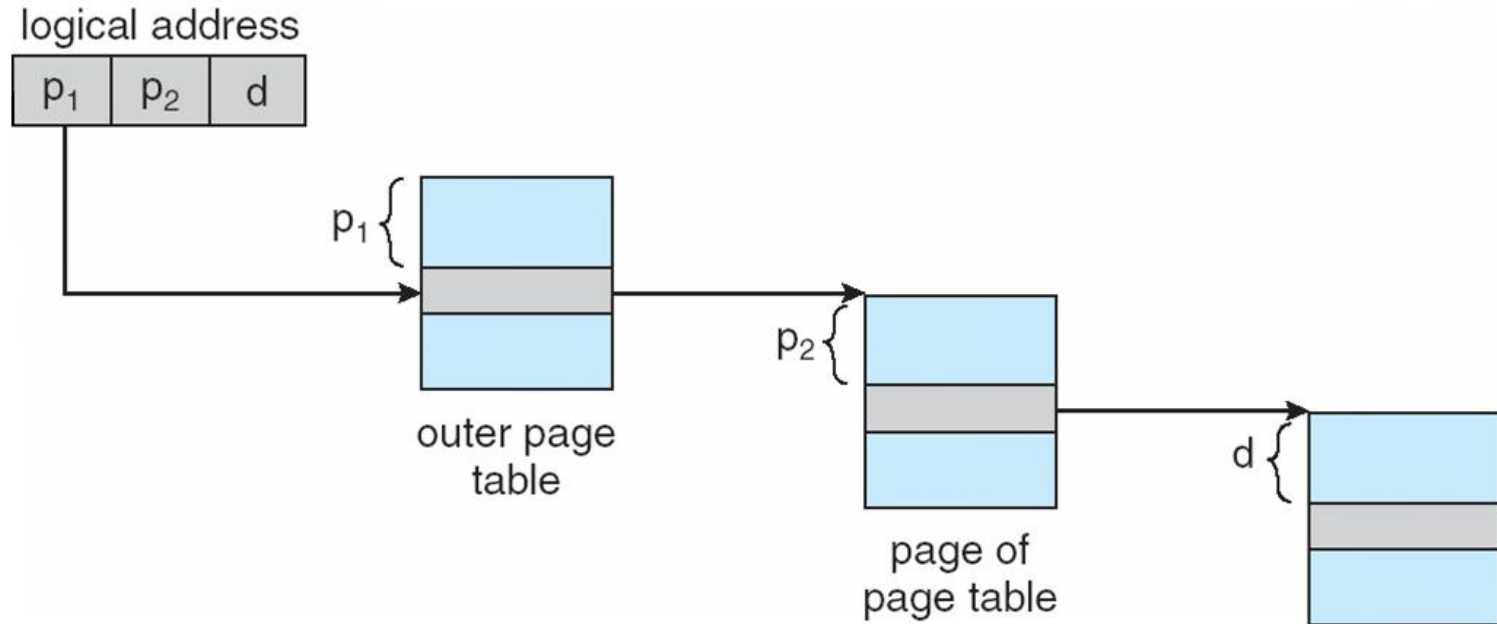
Two-Level Paging Example

- ❏ A logical address (on 32-bit machine with 4K page size) is divided into:
 - ❏ a page number consisting of 20 bits
 - ❏ a page offset consisting of 12 bits
- ❏ The page table should be paged.
- ❏ The page number is further divided into:
 - ❏ a 10-bit page number
 - ❏ a 10-bit page offset
- ❏ Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d
10	10	12

where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

Address-Translation Scheme



**Example: 2-
layer Page Table**

Process Logical address space

Page 0
Page 1
Page 2
Page 3
Page 4
Page 5
Page 6
Page 7
Page 8
Page 9
Page 10
Page 11
Page 12
Page 13
Page 14
Page 15

Page size
=
Frame
size = 8
bytes



Main Memory

	Page
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
...	
$2^{16}-1$	

Example: 2-layer Page Table

Outer Page Table

outer Page #	
0	
1	
2	
3	

Inner Page Table 0

Inner Page #	frame#
0	3
1	7
2	...
3	...

Inner Page Table 1

Inner Page #	frame#
0	...
1	11
2	...
3	0

Inner Page Table 2

Inner Page Table 3

	Page
0	7
1	
2	
3	0
4	
5	
6	
7	1
8	
9	
10	
11	5
12	
13	
...	
$2^{16}-1$	

Main Memory

Page size = 8 bytes

Logical address: 0001001

00: outer-page #
01: inner-page #
001: page offset

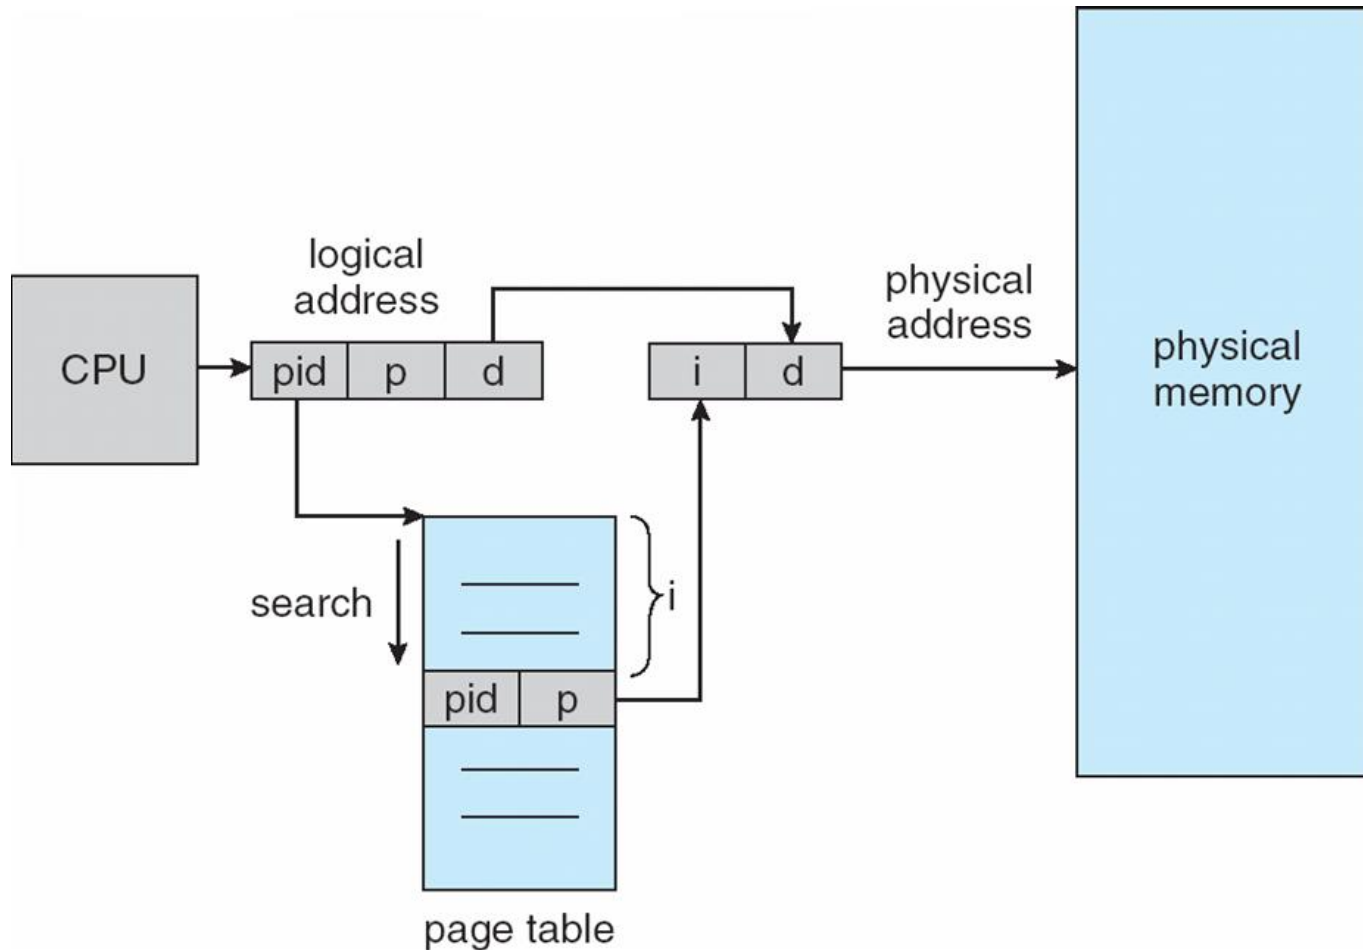
$$\text{frame\#} = (7)_{10} = (111)_2$$

Physical address: 111001

Inverted Page Table

- ❏ One entry for each **frame** of physical memory
- ❏ Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- ❏ Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- ❏ (Use hash table to limit the search to one — or at most a few — page-table entries)

Inverted Page Table Architecture



Process 1's Logical address space

Page 0
Page 1
Page 2
Page 3
Page 4

Process 2's Logical address space

Page 0
Page 1
Page 2
Page 3

Proc 1's Page Table

	frame #
0	7
1	5
2	10
3	12
4	0

Proc 2's Page Table

	frame #
0	1
1	3
2	15
3	8

Main Memory

	Page
0	P1:Page4
1	P2:Page0
2	P3:Page 0
3	P2:Page1
4	P4:Page1
5	P1:Page1
6	P4:Page0
7	P1:Page0
8	P2:Page3
9	P3:Page1
10	P1:Page2
11	P5:Page1
12	P1:Page3
13	P5:Page0
14	P5:Page2
15	P2:Page2

Process 1's Logical address space

Page 0
Page 1
Page 2
Page 3
Page 4

Process 2's Logical address space

Page 0
Page 1
Page 2
Page 3

Inverted Page Table

Frame #	Proc #	Page #
0	1	4
1	2	0
2	3	0
3	2	1
4	4	1
5	1	1
6	4	0
7	1	0
8	2	3
9	3	1
10	1	2
11	5	1
12	1	3
13	5	0
14	5	2
15	2	2

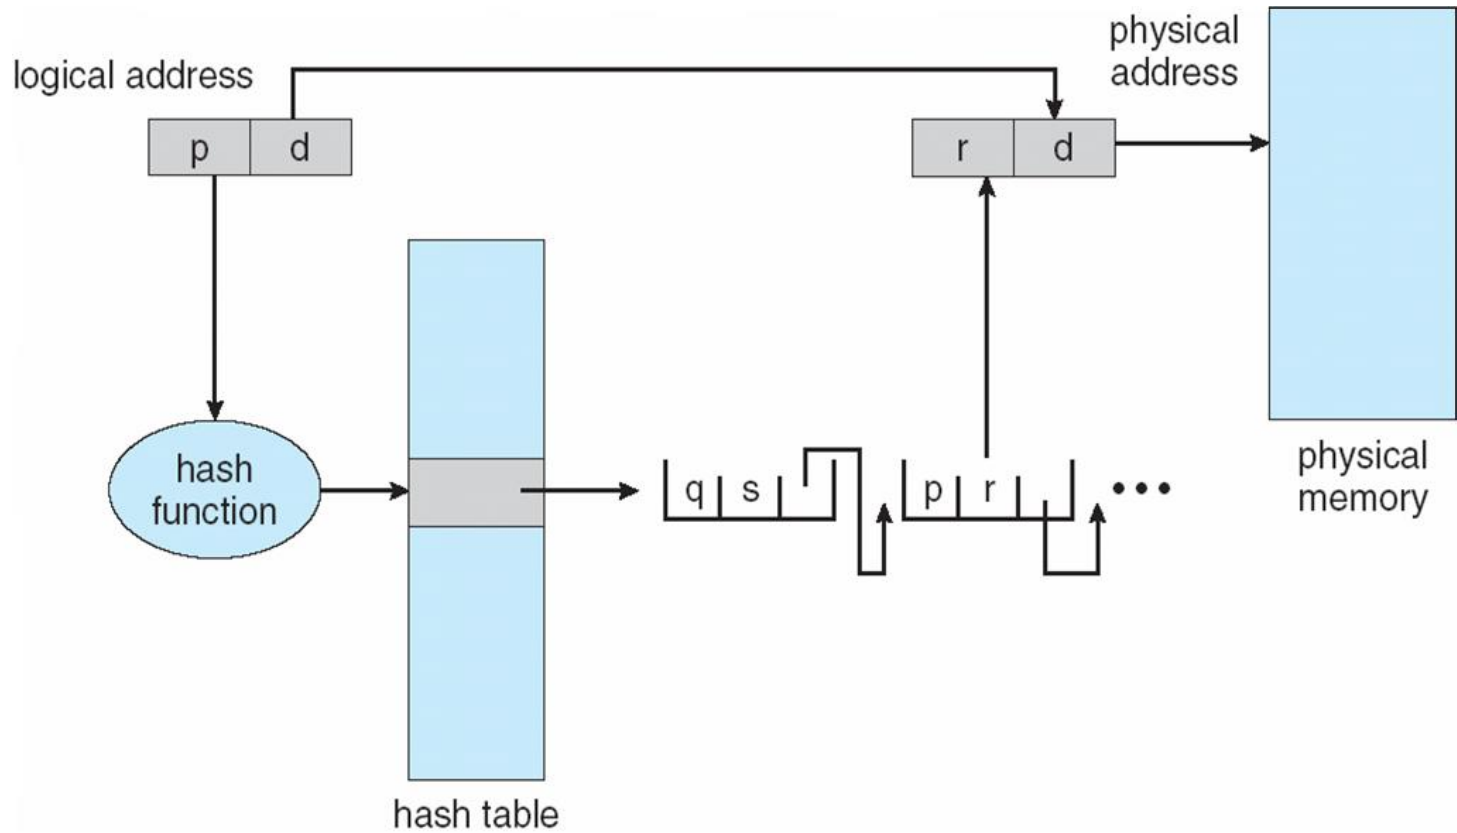
Main Memory

	Page
0	P1:Page4
1	P2:Page0
2	P3:Page 0
3	P2:Page1
4	P4:Page1
5	P1:Page1
6	P4:Page0
7	P1:Page0
8	P2:Page3
9	P3:Page1
10	P1:Page2
11	P5:Page1
12	P1:Page3
13	P5:Page0
14	P5:Page2
15	P2:Page2

Hashed Inverted Page Tables

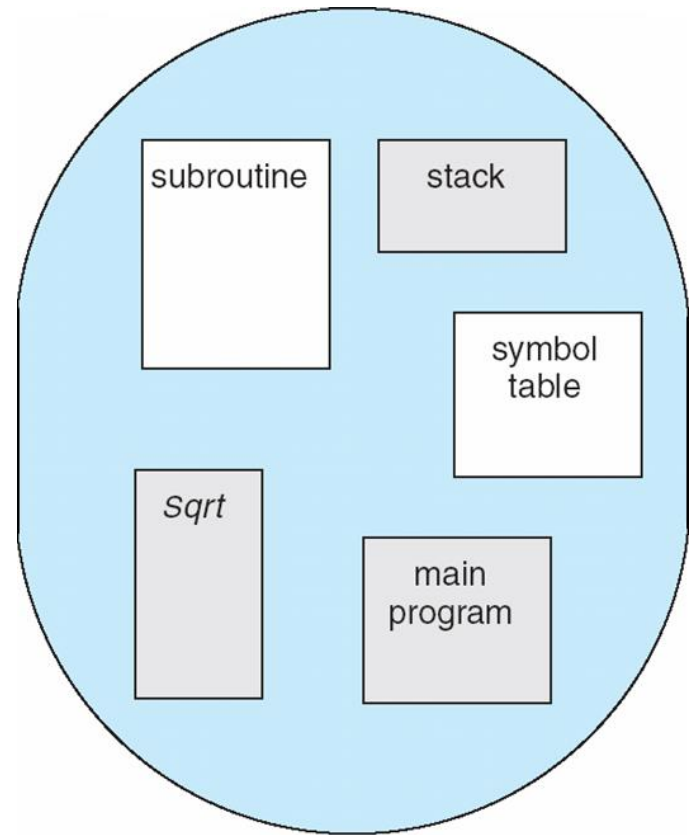
- ❏ The virtual page number is hashed into a table
 - ❏ Each entry of this table contains a chain of elements hashing to the same location
- ❏ Virtual page numbers are compared in this chain searching for a match
 - ❏ If a match is found, the corresponding physical frame is extracted

Hashed Page Table




Segmentation

- ❏ Memory-management scheme that supports user view of memory
- ❏ A program is a collection of segments
 - ❏ A segment is a logical unit such as:
 - main program,
 - procedure/function/method,
 - object,
 - common block,
 - stack (local variables),
 - heap (global variables),
 - ...




logical address

Logical View of Segmentation



user space

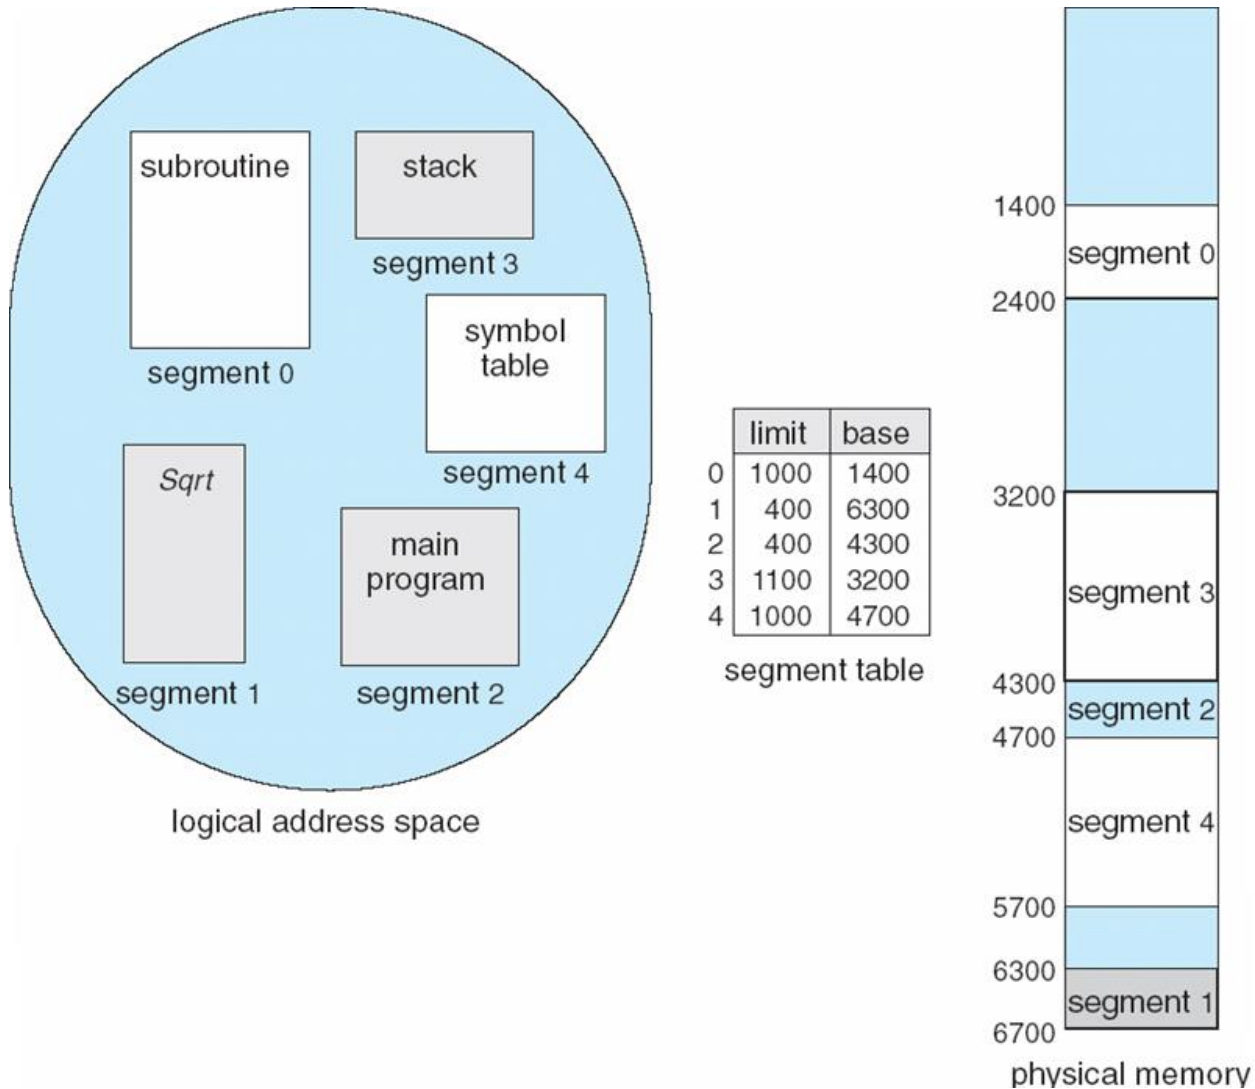


physical memory space

Segmentation Architecture

- ❏ **Segment table** – maps segments to memory; each entry has:
 - ❏ **base** – contains the starting physical address where the segments reside in memory
 - ❏ **limit** – specifies the length of the segment
- ❏ **Segment-table base register (STBR)** points to the segment table's location in memory
- ❏ **Segment-table length register (STLR)** indicates number of segments used by a program

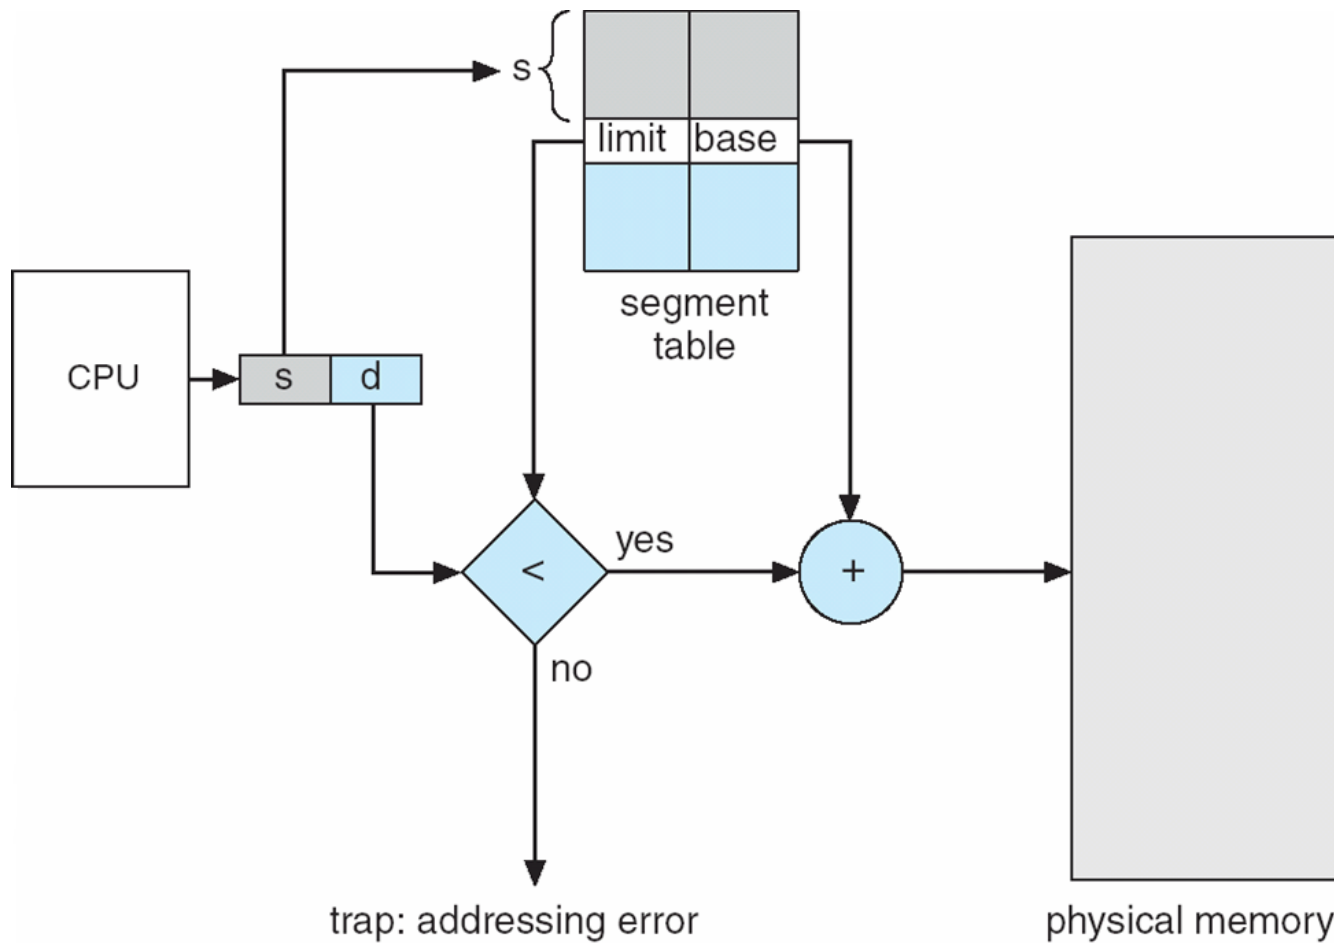
Example of Segmentation



Segmentation Architecture

- Logical address consists of a two tuple: $\langle \text{segment-number}, \text{offset} \rangle$
- Segment number s is legal if $s < \text{STLR}$

Segmentation Hardware





Segmentation Architecture


Protection


 With each entry in segment table associate:

 read/write/execute privileges

 validation bit = 0 \Rightarrow illegal segment

 Protection bits associated with segments; code sharing occurs at segment level

 Since segments vary in length, memory allocation is a dynamic storage-allocation problem

 Could have external fragmentation problem: To address the problem, combination of segmentation and paging may be applied