


Virtual Memory I

October 11, 2017

Virtual Memory: Benefits

- ❏ Only part of a process needs to be in memory for execution
 - ❏ Part of the image of a running process can be in the back store (disk)
- ❏ Logical address space can therefore be much larger than physical address space
 - ❏ It is possible to run a 4G program on a computer with 1G memory
- ❏ Allows for more efficient process creation
 - ❏ Do not need to load in full image




Implementing Virtual Memory: Demand Paging

Key idea:

Bring a page into memory only when it is needed

Pager:

Swapper that deals with pages is also called **pager**




Page Table: Valid-Invalid Bit

- With each page table entry a valid–invalid bit is associated
(**v** \Rightarrow in-memory, **i** \Rightarrow not-in-memory)
- Initially valid–invalid bit is set to **i** on all entries


Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table


Page Table When Some Pages Are Not in Main Memory



Page Fault

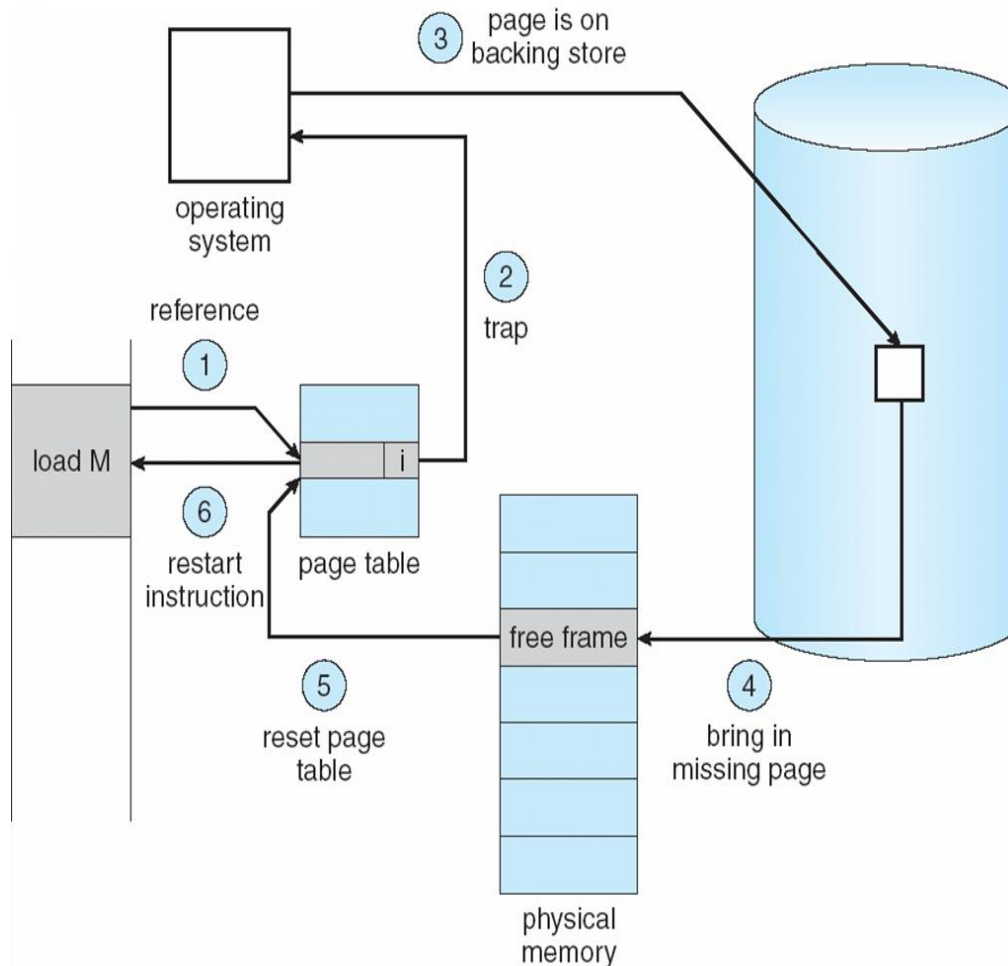
 First reference to a page that is not in memory will trap to operating system: **page fault**

Steps in Handling a Page Fault



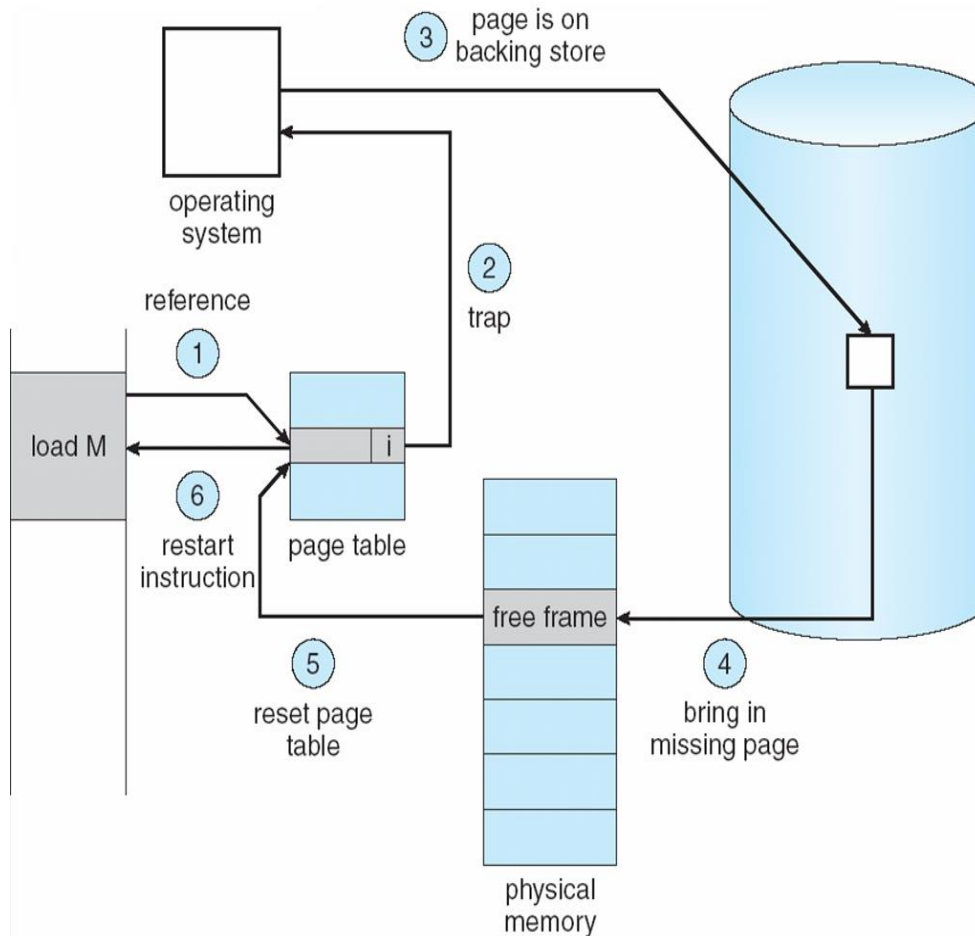
- 1&2. Operating system determine: If invalid reference → abort; if just not in memory → Proceed
3. Identify the desired page in the back store. Get an empty frame in the physical memory (if no empty, swap out one)

Steps in Handling a Page Fault



4. Swap page into frame:
schedule a disk operation to read the desired page into the newly allocated frame. (While waiting for I/O operation to complete, CPU may be allocated to another process; in this case, after I/O operation completes, this process should wait for the CPU to be allocated to it again)

Steps in Handling a Page Fault



5. Modify the page table.
Set validation bit of the newly swapped-in page to **v**
6. Restart the instruction that caused the page fault

Performance of Demand Paging

❏ Page Fault Rate $0 \leq p \leq 1.0$

❏ if $p = 0$, no page faults

❏ if $p = 1$, every reference is a fault

❏ Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) * \text{memory access} \\ & + p * (\text{page fault overhead:} \\ & \quad \text{page fault interrupt service} \\ & \quad + \text{swap page in} \\ & \quad + \text{resume interrupted process} \\ & \quad + \text{memory access} \\ &) \end{aligned}$$



Demand Paging Example

- ❏ Memory access time = 200 nanoseconds
- ❏ Average page-fault overhead = 8 milliseconds
- ❏
$$\begin{aligned} \text{EAT} &= (1 - p) * 200 + p * (8 \text{ milliseconds}) \\ &= (1 - p) * 200 + p * 8,000,000 \\ &= 200 + p * 7,999,800 \end{aligned}$$
- ❏ If one access out of 1,000 causes a page fault, then
EAT = 8.2 microseconds.
This is a slowdown by a factor of 40!!
- ❏ To improve performance:
 - ❏ Reduce page-fault rate: predict pages to access; not swap out pages that will be accessed soon; ...
 - ❏ Reduce page-fault processing time: faster I/O operations

Page Fault Handling: Basic Framework

1. Find the desired page on disk (backing store)
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim** frame
3. Bring the desired page into the free frame; update the page tables
4. Restart the process

Page Replacement

-  Requirement: an ideal page replacement algorithm should result in minimum number of page faults
-  When a read-only page is to be swapped (replaced), do not need to write it back.

FIFO Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2																	
	0	0	0																	
		1	1																	

2	2	4	4	4	0															
3	3	3	2	2	2															
1	0	0	0	3	3															

0	0																			
1	1																			
3	2																			

7	7	7																		
1	0	0																		
2	2	1																		

page frames

First-In-First-Out (FIFO) Algorithm

Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

3 frames (3 pages can be in memory at a time per process)

1	1	4	5	
2	2	1	3	9 page faults
3	3	2	4	

4 frames

1	1	5	4	
2	2	1	5	10 page faults
3	3	2		
4	4	3		

Belady's Anomaly: more frames \Rightarrow more page faults

Optimal (Ideal) Algorithm

❏ Replace page that will not be used for the longest period of time

❏ 4 frames example

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1
2
3
4

4

6 page faults

5


❏ How do you know this? No.

❏ Why do we care this algorithm? Used for measuring how well your algorithm performs

Optimal Page Replacement

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames

Least Recently Used (LRU) Algorithm

 Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

1	1	1	1	5
2	2	2	2	2
3	5	5	4	4
4	4	3	3	3

LRU Page Replacement

Counter implementation

- Every frame has a counter; every time the page in a frame is referenced, copy the clock into the counter of the frame
- When a page needs to be replaced, the page in the frame with the oldest counter value is to be replaced

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0			1		1		1	
	0	0	0		0		0	0	3	3			3		0		0	
		1	1		3		3	2	2	2			2		2		7	

page frames


LRU Algorithm: Stack Implementation

- Keep a stack of page numbers
- Page referenced is in the stack: move it to the top
- Page referenced is not in the stack:
 - If there is free frame → push the page into the stack
 - If there is no free frame → swap out the page on the stack bottom; push the new page into the stack
- Example
 - Reference string 70120304230321201701
 - #of frames=3


Second-Chance: An Approximate LRU

- ❏ A reference bit associated with each page in physical memory
- ❏ All pages in physical memory form a circular queue; a pointer pointing to the head element
- ❏ When a page is referenced, its reference bit is set to 1.
- ❏ When a page should be replaced:
 - ❏ Step 1. Move the pointer by one step
 - ❏ Step 2. Check the page pointed by the pointer:
 - ❏ If the associated bit is 0 → replace the page
 - ❏ If the associated bit is 1 → change the bit to 0 and go to step 1.

Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)


Minimal Number of Pages Needed

- Each process needs *minimum* number of pages: determined by computer architecture
- Example: IBM 370 needs 6 pages to handle MOVE instruction:
 - instruction is 6 bytes, might span 2 pages
 - 2 pages to handle *from*
 - 2 pages to handle *to*

How many pages are “enough”?


Thrashing

- ❏ If a process does not have “enough” pages, the page-fault rate is very high. This leads to:
 - ❏ low CPU utilization: handling page-fault; frequent proc scheduling
 - ❏ (because CPU is not fully used) OS thinks that it needs to increase the degree of multiprogramming: another process added to the system
- ❏ **Thrashing** \equiv a process is busy swapping pages in and out



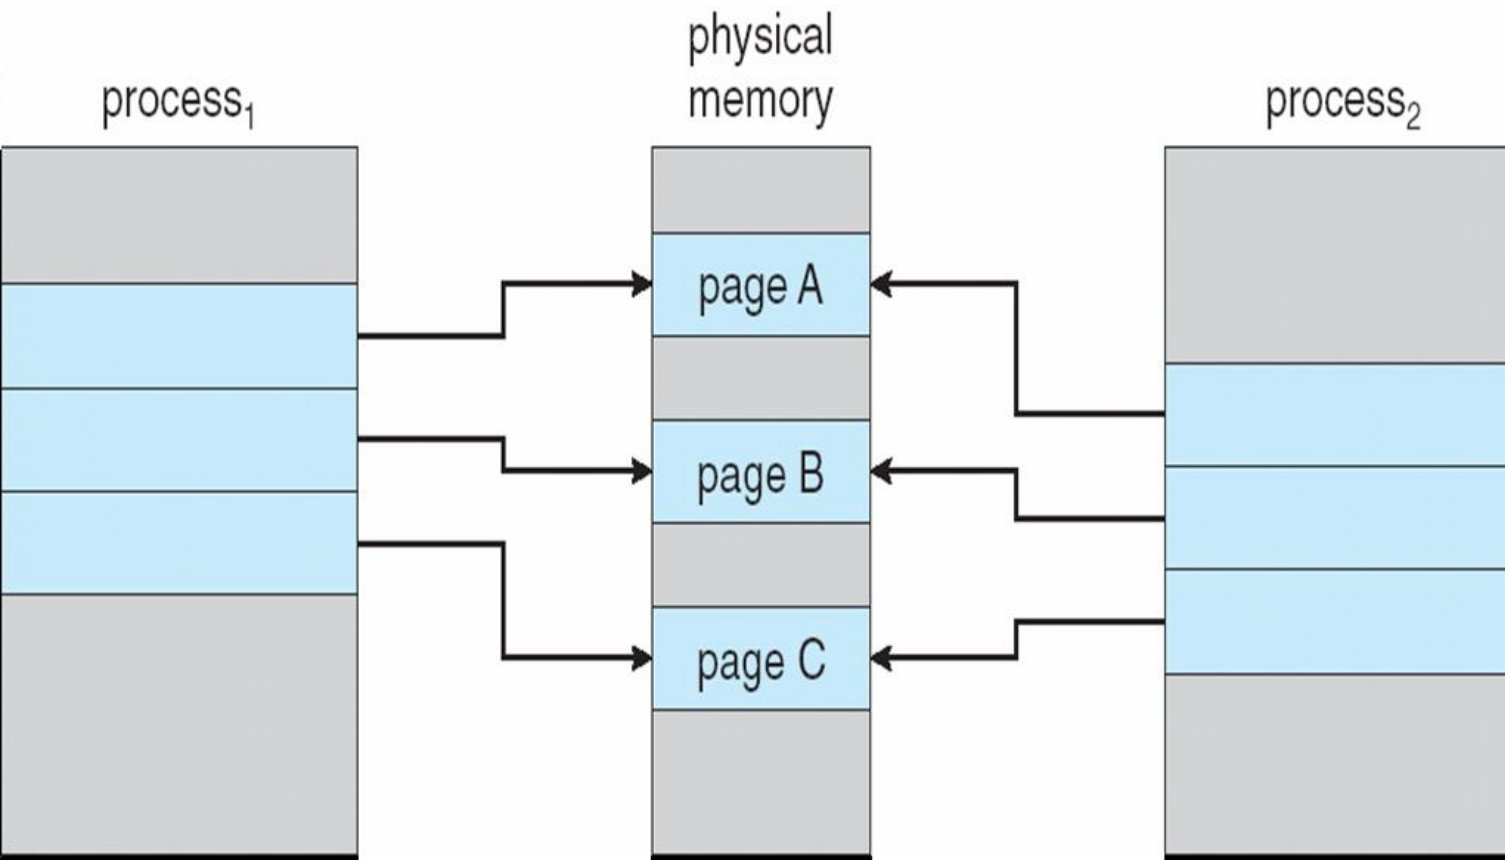
Page-Fault Frequency Scheme

- ❏ Establish “acceptable” page-fault rate for a system
- ❏ Keep track of the actual page-fault rate for each process
 - ❏ If actual rate too low, process loses frame
 - ❏ If actual rate too high, process gains frame



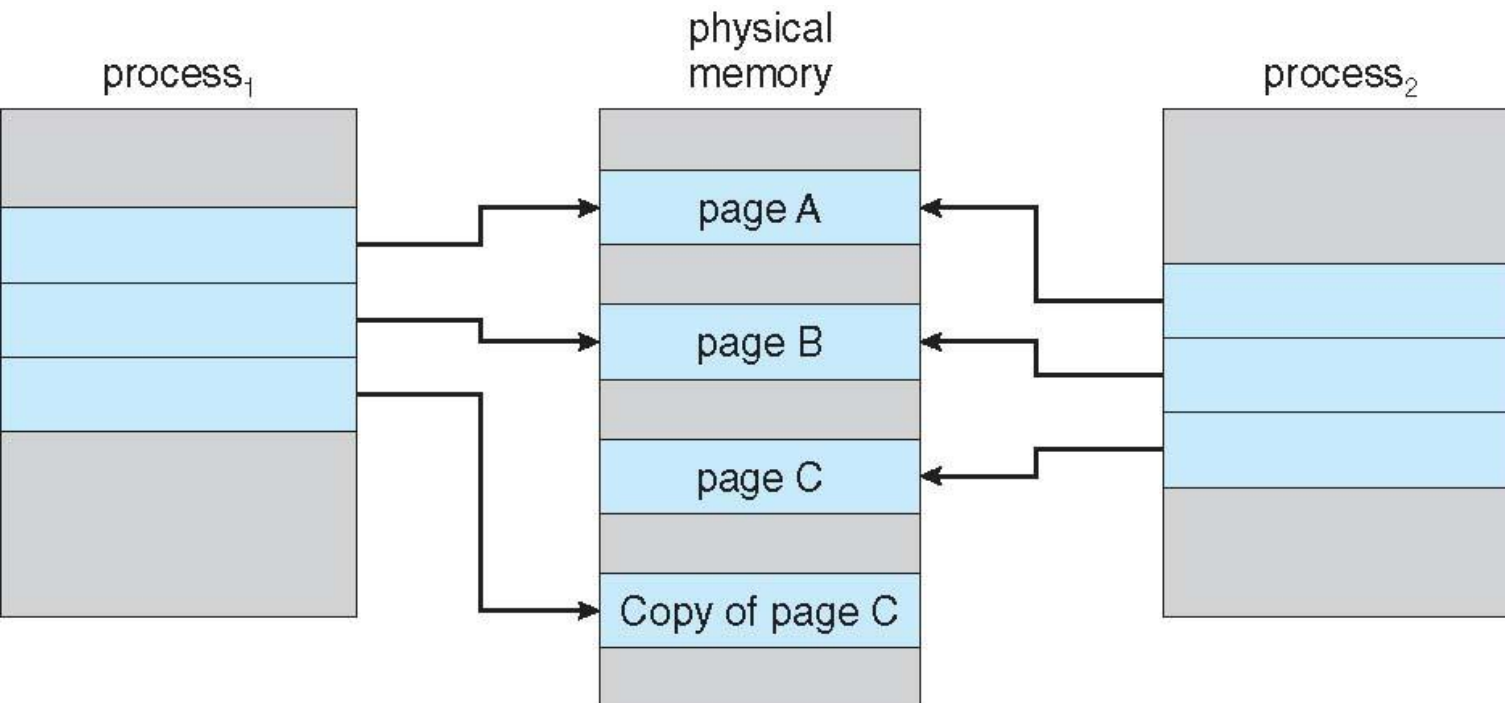
Copy-on-Write: Speed Up Process Creation

- When a new process is created, it copies the image of its parent.
- Copy-on-Write (COW) allows both parent and child processes to initially *share* the same pages in memory



Copy-on-Write: Speed Up Process Creation

- If either process modifies a shared page, only then is the page copied



Memory-Mapped Files

- ❏ Memory-mapped file I/O allows file I/O to be treated as routine memory access by **mapping** a disk block to a page in memory
- ❏ A file is initially read using demand paging. A page-sized portion of the file is read from the file system into a physical page. Subsequent reads/writes to/from the file are treated as ordinary memory accesses.
- ❏ Simplifies file access by treating file I/O through memory rather than **read()** **write()** system calls.

Memory Mapped Files

